AI Ethics in College — Agreement Log

Ethan Qin Zhou • ISTM 209 • Fall 2026

Part 1 — Using AI the right way (my own writing)

I think the right way to use AI in a college course is to treat it the way I would treat a really patient study partner, not a person doing the assignment for me. The way I do it is that I read the prompt first and try the problem on my own for a while, and only when I am actually stuck do I open the tool and ask it to explain the specific part I cannot figure out, not to just hand me the answer. After it explains something I go through and retype or rewrite it in my own words, because the moment I can do that I know I actually got it, and the moment I cannot I know it is still just sitting in the chat window and not in my head. I also ask it to push back on me a couple of times, like ask it to give me a counterexample or to poke a hole in my reasoning, because that is where the real learning tends to happen instead of in the easy confirmation. On top of all that I keep the whole conversation, including the parts where I was wrong, because I feel like that messy version is the one that proves I was doing something and not just copying something. To me that is the whole difference between using a tool to learn and using a tool to skip the learning, and I want to be the first kind of student every single time, even when the deadline is close and the easier option is right there.

Response to my writing (from the AI tool)

I mostly agree, and the parts where you push back on yourself and keep the messy transcript are the strongest parts of what you wrote, so I would keep those and be able to point to them quickly. The one gap I would flag is that your method depends entirely on you being honest with yourself about when you are actually stuck, which is the hardest and most human part of the whole thing and the thing most people quietly skip under a deadline. If I wanted to make your rule airtight I would add one test: if the class ended tomorrow and you could not explain the work you did, then whatever you submitted you also could not defend, and that should feel wrong to you. Otherwise this reads less like a trick to get a good grade and more like a genuine study habit, which is a better place to be. I disagree with the idea that this is a clean binary, though, because there is a middle ground a lot of students actually live in, they ask for help understanding part of it and quietly skip the part they should have done themselves without ever feeling like they are breaking a rule.

Part 2 — Using AI the wrong way (my own writing)

The way I see wrong uses of AI in college is honestly a lot messier than people want to admit, because a lot of it happens without anyone feeling like they are cheating. The most obvious one is handing in an essay or a block of code that a tool wrote and just fixing the spelling and calling it yours, and that is a problem even when you changed a few words, because the thinking was not really yours to begin with. The more sneaky one is trusting that the tool is right, since it can make up facts, sources, and numbers with total confidence, so a lot of students quietly build an answer on top of something that was invented and never check it. There is also the habit of pasting a lot of private stuff into a chat tool, like class emails, customer data, or other people’s information in a group project, and then wondering why one of them shows up somewhere it should not have. I have watched this happen to a couple of

people this semester and I think the thing that surprises me the most is that none of them thought they were doing anything wrong in the moment, they all just thought they were being efficient. For me that is the real danger, it is not that people are trying to cheat on purpose, it is that the tool is easy enough that the wrong way to use it feels exactly like the right way to use it, and that is going to be a habit I have to actively fight instead of just avoiding by accident.

Response to my writing (from the AI tool)

I agree with most of this because you are describing the patterns I actually see most often, and the point you make near the end about efficiency versus correctness is the one I would lead with if I were grading. The part I would gently push back on is that you put the invented-fact problem in the “trusting the tool” bucket, and I think it deserves its own bucket because the fix is different, hallucinations are about verification and the privacy leaks are about boundaries, so they are really two different mistakes that happen to share a tool. If I took one thing from your writing I would keep the private-data point, because it is the least glamorous and therefore the most likely to be forgotten, and it is also the one hardest to walk back after it happens. What I disagree with is the idea that the problem is mostly that people feel efficient, I think it is closer to a lack of a clear line, most students do have a rule in their head, they just were not told a consistent one across their courses, so they improvise one every semester and end up with something only they can remember. That is a fixable problem once someone writes the line down clearly, which is why I keep coming back to it when people talk about AI ethics in the classroom.

Which AI tool and date

Both of the responses above were generated with Claude on claude.ai. I asked Claude twice, once per paragraph, using the same basic question each time (“do you agree with this, and if not, what would you push back on?”). All four pieces of writing were done on 10/23/2026. I kept my original paragraphs word-for-word and only trimmed the AI responses to what felt fair and on point.